

vjy3u9nei

January 25, 2025

1 Assignment 1

Question 1: [3 Marks]

Explain the functionality of a perceptron with its mathematical representation. Provide:

The formula for the perceptron model.

A detailed explanation of each term in the formula.

Question 2: [4 Marks]

Build and train a neural network to solve the XOR problem using any deep learning library (e.g., TensorFlow or PyTorch):

Define the 3-bit XOR dataset and preprocess it for model training. [1 Mark]

Design a neural network with one hidden layer. Clearly specify the architecture: [input, hidden, output]. [2 Marks] Train the network on the dataset and evaluate its performance, reporting the accuracy. [1 Mark]

Question 3: [5 Marks]

Implement a neural network for the Pima Indians Diabetes Dataset:

Design a neural network with the following architecture: [Input, hidden1(8), hidden2(4), hidden3(4), output]. [2 Marks]

Train the model on the diabetes dataset and record its performance. [2 Marks]

Provide a schematic representation of the neural network with its layers and activation functions. [1 Mark]

1.1 Question 1: [3 Marks]

Explain the functionality of a perceptron with its mathematical representation. Provide:

1 The formula for the perceptron model

2. A detailed explanation of each term in the formula.

1.1.1 1- The formula for the perceptron model.

A perceptron is a simple artificial neuron designed for binary classification. It maps inputs to an output based on a weighted sum and a threshold. Below is the explanation of its functionality, formula, and the detailed explanation of each term. The perceptron is a fundamental building block

of neural networks used for binary classification. It processes inputs and computes an output based on a linear combination of weights, inputs, and a bias term, followed by an activation function.

The perceptron works on the following principle:

- > Takes weighted sums of input features.
- > Passes the result through an activation function (typically a step function).
- > Outputs a binary decision (0 or 1).

Mathematical Representation of a Perceptron The perceptron model can be described using the following formula:

$$y = f\left(\sum_{i=1}^n w_i x_i + b\right)$$

Where:

- > y : The output of the perceptron, which is either 0 or 1.
- > $f(z)$: The activation function, which applies a step function
- > $w_i = [w_1, w_2, \dots, w_n]$: The weights of the model associated with i -th input feature.
- > $x_i = [x_1, x_2, \dots, x_n]$: i -th input features.
- > b : The bias term, which shifts the decision boundary.
- > $\sum_{i=1}^n w_i x_i$: Weighted sum of inputs.

1.1.2 2- A detailed explanation of each term in the formula.

1. Input Features (x_i):

These are the feature values of the input data. For example, if the input is a vector like $[1, 2, 3]$ each represents a specific attribute or feature.

2. Weights (w_i):

Weights represent the importance or contribution of each input feature to the output. The weights are adjusted during training to minimize classification errors.

3. Weighted Sum ($\sum_{i=1}^n w_i x_i$):

This is the linear combination of inputs and their corresponding weights. It determines the overall influence of the inputs before applying the activation function.

4. Bias (b):

The bias term shifts the decision boundary away from the origin, allowing more flexibility in classification. Without bias, the decision boundary always passes through the origin.

5. Activation Function ($f(z)$):

The activation function determines the output of the perceptron. In a simple perceptron, the step function is used:

$$f(z) = \begin{cases} 1, & \text{if } z \geq 0 \\ 0, & \text{otherwise} \end{cases}$$

$$0, \text{if } < 0$$

Here, $= \sum_{i=1}^n w_i x_i + b$)

b is the input to the activation function.

Summary:-

The perceptron essentially calculates a weighted sum of inputs, adjusts it using a bias, and applies an activation function to make a classification decision. Its functionality can be visualized geometrically as finding a hyperplane (decision boundary) that separates the data into two classes in a feature space.

1.2 Question 2: [4 Marks]

Build and train a neural network to solve the XOR problem using any deep learning library (e.g., TensorFlow or PyTorch):

1. Define the 3-bit XOR dataset and preprocess it for model training. [1 Mark]
2. Design a neural network with one hidden layer. Clearly specify the architecture: [input, hidden, output]. [2 Marks]
3. Train the network on the dataset and evaluate its performance, reporting the accuracy. [1 Mark]

1.2.1 1. Define the 3-bit XOR dataset and preprocess it for model training. [1 Mark]

```
[3]: # Importing required packages
import numpy as np
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense

# Defining 3-bit XOR dataset
x = np.array([
    [0, 0, 0],
    [0, 0, 1],
    [0, 1, 0],
    [0, 1, 1],
    [1, 0, 0],
    [1, 0, 1],
    [1, 1, 0],
    [1, 1, 1]
])

# XOR output (1 if odd number of 1s, otherwise 0)
y = np.array([[0], [1], [1], [0], [1], [0], [0], [1]])

# Normalize the inputs (not strictly necessary here since inputs are binary)
x = x.astype("float32")
```

```
y = y.astype("float32")
```

```
[4]: x
```

```
[4]: array([[0., 0., 0.],
          [0., 0., 1.],
          [0., 1., 0.],
          [0., 1., 1.],
          [1., 0., 0.],
          [1., 0., 1.],
          [1., 1., 0.],
          [1., 1., 1.]], dtype=float32)
```

```
[5]: y
```

```
[5]: array([[0.],
          [1.],
          [1.],
          [0.],
          [1.],
          [0.],
          [0.],
          [1.]], dtype=float32)
```

1.2.2 2. Design a neural network with one hidden layer. Clearly specify the architecture: [input, hidden, output]. [2 Marks]

Designing a neural network with one hidden layer:

- Input layer: 3 neurons (corresponding to the 3 input bits).
- Hidden layer: 4 neurons with ReLU activation.
- Output layer: 1 neuron with sigmoid activation (for binary classification).

```
[6]: # Designing the neural network
model = Sequential([
    Dense(4, activation="relu", input_shape=(3,)), # Hidden layer with 4
    ↪neurons
    Dense(1, activation="sigmoid")                 # Output layer with 1 neuron
])

# Compile the model
model.compile(optimizer="adam", loss="binary_crossentropy",
    ↪metrics=["accuracy"])
```

C:\Users\ASUS\miniconda3\Lib\site-packages\keras\src\layers\core\dense.py:87:
UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When
using Sequential models, prefer using an `Input(shape)` object as the first

layer in the model instead.

```
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

1.2.3 3. Train the network on the dataset and evaluate its performance, reporting the accuracy. [1 Mark]

```
[9]: # Training the model
model.fit(x, y, epochs=100, verbose=0, batch_size=8)

# Evaluate the model
loss, accuracy = model.evaluate(x, y, verbose=0)

print(f"Final Model Accuracy: {accuracy * 100:.2f}%")
```

Final Model Accuracy: 75.00%

1.3 Question 3: [5 Marks]

Implement a neural network for the Pima Indians Diabetes Dataset: 1. Design a neural network with the following architecture: [Input, hidden1(8), hidden2(4), hidden3(4), output]. [2 Marks] 2. Train the model on the diabetes dataset and record its performance. [2 Marks] 3. Provide a schematic representation of the neural network with its layers and activation functions. [1 Mark]

1.3.1 1. Design a neural network with the following architecture: [Input, hidden1(8), hidden2(4), hidden3(4), output]. [2 Marks]

```
[10]: # Importing required packages
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
import warnings as war
war.filterwarnings("ignore")
```

```
[11]: # Step 1:- Load the Diabetes Dataset

dataSetPath="C:\\Users\\ASUS\\jupyterworkspace\\Assignment & Mini_
↳Project\\Module_06_Deep_
↳Learning\\Deep-Learning-Assignment01-neuralnetworkdesignon_diabetesdataSet\\diabetes.
↳csv"
dataSetRead=pd.read_csv(dataSetPath)
```

```
[12]: # Displaying first 5 records to confirming data loading
print("*****Displaying below_
↳first 5 records*****")
dataSetRead.head()
```

```
*****Displaying below first 5
records*****
```

```
[12]: Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin   BMI   \
0           6        148             72             35         0  33.6
1           1         85             66             29         0  26.6
2           8        183             64              0         0  23.3
3           1         89             66             23        94  28.1
4           0        137             40             35       168  43.1

DiabetesPedigreeFunction  Age  Outcome
0                0.627    50         1
1                0.351    31         0
2                0.672    32         1
3                0.167    21         0
4                2.288    33         1
```

```
[13]: # Displaying last 5 records to confirming data loading
print("*****Displaying below
↳last 5 records*****")
dataSetRead.tail()
```

```
*****Displaying below last 5
records*****
```

```
[13]: Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin   BMI   \
763          10        101             76             48       180  32.9
764           2        122             70             27         0  36.8
765           5        121             72             23       112  26.2
766           1        126             60              0         0  30.1
767           1         93             70             31         0  30.4

DiabetesPedigreeFunction  Age  Outcome
763                0.171    63         0
764                0.340    27         0
765                0.245    30         0
766                0.349    47         1
767                0.315    23         0
```

```
[14]: # Displaying all records to confirming data loading
print("*****Displaying below
↳all records*****")
dataSetRead
```

```
*****Displaying below all
records*****
```

```
[14]: Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin   BMI   \
0           6        148             72             35         0  33.6
1           1         85             66             29         0  26.6
```

2	8	183	64	0	0	23.3
3	1	89	66	23	94	28.1
4	0	137	40	35	168	43.1
..
763	10	101	76	48	180	32.9
764	2	122	70	27	0	36.8
765	5	121	72	23	112	26.2
766	1	126	60	0	0	30.1
767	1	93	70	31	0	30.4

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1
..
763	0.171	63	0
764	0.340	27	0
765	0.245	30	0
766	0.349	47	1
767	0.315	23	0

[768 rows x 9 columns]

```
[15]: # Splitting data into features (X) and target (y)
X = dataSetRead.iloc[:, :-1].values # All columns except the last one
y = dataSetRead.iloc[:, -1].values # Last column as the target variable

# Splitting the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳ random_state=42)

# Standardizing the features (important for neural networks)
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

1.3.2 Design the Neural Network Architecture

The neural network will have the following structure:

- > Input Layer: Matches the number of features in the dataset (8 features).
- > Hidden Layer 1: 8 neurons, ReLU activation.
- > Hidden Layer 2: 4 neurons, ReLU activation.
- > Hidden Layer 3: 4 neurons, ReLU activation.

→ Output Layer: 1 neuron, sigmoid activation (for binary classification).

```
[16]: # Defining the model
model = Sequential([
    Dense(8, activation='relu', input_shape=(X.shape[1],)), # Hidden layer 1
    Dense(4, activation='relu'), # Hidden layer 2
    Dense(4, activation='relu'), # Hidden layer 3
    Dense(1, activation='sigmoid') # Output layer
])

# Compiling the model
model.compile(optimizer='adam', loss='binary_crossentropy',
    ↪metrics=['accuracy'])
```

1.3.3 2. Train the model on the diabetes dataset and record its performance. [2 Marks]

```
[19]: # Training the model
history = model.fit(X_train, y_train, epochs=100, batch_size=16,
    ↪validation_split=0.2, verbose=1)
```

```
Epoch 1/100
31/31          0s 14ms/step -
accuracy: 0.8178 - loss: 0.4344 - val_accuracy: 0.7154 - val_loss: 0.5904
Epoch 2/100
31/31          0s 11ms/step -
accuracy: 0.8114 - loss: 0.4465 - val_accuracy: 0.7073 - val_loss: 0.5918
Epoch 3/100
31/31          0s 11ms/step -
accuracy: 0.8166 - loss: 0.4365 - val_accuracy: 0.7154 - val_loss: 0.5891
Epoch 4/100
31/31          0s 12ms/step -
accuracy: 0.7823 - loss: 0.4601 - val_accuracy: 0.7073 - val_loss: 0.5904
Epoch 5/100
31/31          0s 12ms/step -
accuracy: 0.7979 - loss: 0.4568 - val_accuracy: 0.7154 - val_loss: 0.5917
Epoch 6/100
31/31          0s 12ms/step -
accuracy: 0.8032 - loss: 0.4651 - val_accuracy: 0.7073 - val_loss: 0.5930
Epoch 7/100
31/31          0s 11ms/step -
accuracy: 0.7863 - loss: 0.4556 - val_accuracy: 0.7073 - val_loss: 0.5923
Epoch 8/100
31/31          0s 11ms/step -
accuracy: 0.8275 - loss: 0.4117 - val_accuracy: 0.7073 - val_loss: 0.5919
Epoch 9/100
31/31          0s 11ms/step -
accuracy: 0.8075 - loss: 0.4422 - val_accuracy: 0.7073 - val_loss: 0.5928
```


Epoch 10/100
31/31 0s 11ms/step -
accuracy: 0.7837 - loss: 0.4869 - val_accuracy: 0.7073 - val_loss: 0.5936
Epoch 11/100
31/31 0s 11ms/step -
accuracy: 0.8096 - loss: 0.4419 - val_accuracy: 0.7073 - val_loss: 0.5919
Epoch 12/100
31/31 0s 11ms/step -
accuracy: 0.7951 - loss: 0.4658 - val_accuracy: 0.7154 - val_loss: 0.5928
Epoch 13/100
31/31 0s 11ms/step -
accuracy: 0.7642 - loss: 0.4896 - val_accuracy: 0.7154 - val_loss: 0.5943
Epoch 14/100
31/31 0s 12ms/step -
accuracy: 0.7934 - loss: 0.4803 - val_accuracy: 0.7236 - val_loss: 0.5955
Epoch 15/100
31/31 0s 11ms/step -
accuracy: 0.8023 - loss: 0.4292 - val_accuracy: 0.7154 - val_loss: 0.5950
Epoch 16/100
31/31 0s 11ms/step -
accuracy: 0.8325 - loss: 0.4128 - val_accuracy: 0.7154 - val_loss: 0.5973
Epoch 17/100
31/31 0s 12ms/step -
accuracy: 0.8117 - loss: 0.4287 - val_accuracy: 0.7154 - val_loss: 0.5936
Epoch 18/100
31/31 0s 13ms/step -
accuracy: 0.8203 - loss: 0.4257 - val_accuracy: 0.7154 - val_loss: 0.5957
Epoch 19/100
31/31 0s 11ms/step -
accuracy: 0.7857 - loss: 0.4610 - val_accuracy: 0.7154 - val_loss: 0.5965
Epoch 20/100
31/31 0s 11ms/step -
accuracy: 0.8051 - loss: 0.4488 - val_accuracy: 0.7073 - val_loss: 0.5946
Epoch 21/100
31/31 0s 11ms/step -
accuracy: 0.8060 - loss: 0.4448 - val_accuracy: 0.7154 - val_loss: 0.5967
Epoch 22/100
31/31 0s 11ms/step -
accuracy: 0.8347 - loss: 0.4212 - val_accuracy: 0.7073 - val_loss: 0.5959
Epoch 23/100
31/31 0s 13ms/step -
accuracy: 0.8089 - loss: 0.4315 - val_accuracy: 0.7154 - val_loss: 0.5994
Epoch 24/100
31/31 0s 11ms/step -
accuracy: 0.8092 - loss: 0.4520 - val_accuracy: 0.7154 - val_loss: 0.5973
Epoch 25/100
31/31 1s 15ms/step -
accuracy: 0.8107 - loss: 0.4414 - val_accuracy: 0.7154 - val_loss: 0.5996

Epoch 26/100
 31/31 0s 13ms/step -
 accuracy: 0.7956 - loss: 0.4661 - val_accuracy: 0.7154 - val_loss: 0.6011
 Epoch 27/100
 31/31 0s 11ms/step -
 accuracy: 0.8024 - loss: 0.4437 - val_accuracy: 0.7154 - val_loss: 0.5984
 Epoch 28/100
 31/31 0s 12ms/step -
 accuracy: 0.7834 - loss: 0.4834 - val_accuracy: 0.7154 - val_loss: 0.6000
 Epoch 29/100
 31/31 0s 12ms/step -
 accuracy: 0.8120 - loss: 0.4428 - val_accuracy: 0.7154 - val_loss: 0.5989
 Epoch 30/100
 31/31 0s 13ms/step -
 accuracy: 0.8250 - loss: 0.4379 - val_accuracy: 0.7154 - val_loss: 0.6008
 Epoch 31/100
 31/31 0s 12ms/step -
 accuracy: 0.8148 - loss: 0.4397 - val_accuracy: 0.7154 - val_loss: 0.5998
 Epoch 32/100
 31/31 0s 11ms/step -
 accuracy: 0.8158 - loss: 0.4360 - val_accuracy: 0.7154 - val_loss: 0.6011
 Epoch 33/100
 31/31 0s 11ms/step -
 accuracy: 0.8303 - loss: 0.4261 - val_accuracy: 0.7154 - val_loss: 0.6027
 Epoch 34/100
 31/31 0s 11ms/step -
 accuracy: 0.7862 - loss: 0.4672 - val_accuracy: 0.7154 - val_loss: 0.6022
 Epoch 35/100
 31/31 0s 11ms/step -
 accuracy: 0.8494 - loss: 0.4047 - val_accuracy: 0.7154 - val_loss: 0.6024
 Epoch 36/100
 31/31 0s 12ms/step -
 accuracy: 0.7958 - loss: 0.4469 - val_accuracy: 0.7154 - val_loss: 0.6038
 Epoch 37/100
 31/31 0s 11ms/step -
 accuracy: 0.8039 - loss: 0.4458 - val_accuracy: 0.7154 - val_loss: 0.6031
 Epoch 38/100
 31/31 0s 12ms/step -
 accuracy: 0.8217 - loss: 0.4128 - val_accuracy: 0.7154 - val_loss: 0.6026
 Epoch 39/100
 31/31 0s 11ms/step -
 accuracy: 0.8128 - loss: 0.4187 - val_accuracy: 0.7154 - val_loss: 0.6052
 Epoch 40/100
 31/31 0s 11ms/step -
 accuracy: 0.8401 - loss: 0.4218 - val_accuracy: 0.7154 - val_loss: 0.6033
 Epoch 41/100
 31/31 0s 11ms/step -
 accuracy: 0.8207 - loss: 0.4191 - val_accuracy: 0.7236 - val_loss: 0.6068

Epoch 42/100
 31/31 0s 12ms/step -
 accuracy: 0.8217 - loss: 0.4147 - val_accuracy: 0.7154 - val_loss: 0.6015
 Epoch 43/100
 31/31 1s 15ms/step -
 accuracy: 0.8174 - loss: 0.4234 - val_accuracy: 0.7154 - val_loss: 0.6034
 Epoch 44/100
 31/31 0s 12ms/step -
 accuracy: 0.8289 - loss: 0.4352 - val_accuracy: 0.7236 - val_loss: 0.6070
 Epoch 45/100
 31/31 0s 11ms/step -
 accuracy: 0.8087 - loss: 0.4399 - val_accuracy: 0.7154 - val_loss: 0.6049
 Epoch 46/100
 31/31 0s 11ms/step -
 accuracy: 0.7805 - loss: 0.4829 - val_accuracy: 0.7154 - val_loss: 0.6039
 Epoch 47/100
 31/31 0s 10ms/step -
 accuracy: 0.8155 - loss: 0.4306 - val_accuracy: 0.7154 - val_loss: 0.6088
 Epoch 48/100
 31/31 0s 11ms/step -
 accuracy: 0.7850 - loss: 0.4766 - val_accuracy: 0.7154 - val_loss: 0.6063
 Epoch 49/100
 31/31 1s 12ms/step -
 accuracy: 0.8020 - loss: 0.4443 - val_accuracy: 0.7236 - val_loss: 0.6055
 Epoch 50/100
 31/31 0s 11ms/step -
 accuracy: 0.8080 - loss: 0.4374 - val_accuracy: 0.7154 - val_loss: 0.6088
 Epoch 51/100
 31/31 0s 11ms/step -
 accuracy: 0.8057 - loss: 0.4460 - val_accuracy: 0.7154 - val_loss: 0.6090
 Epoch 52/100
 31/31 0s 12ms/step -
 accuracy: 0.8132 - loss: 0.4368 - val_accuracy: 0.7236 - val_loss: 0.6076
 Epoch 53/100
 31/31 0s 11ms/step -
 accuracy: 0.7809 - loss: 0.4558 - val_accuracy: 0.7154 - val_loss: 0.6098
 Epoch 54/100
 31/31 0s 13ms/step -
 accuracy: 0.8152 - loss: 0.4323 - val_accuracy: 0.7236 - val_loss: 0.6092
 Epoch 55/100
 31/31 0s 11ms/step -
 accuracy: 0.8138 - loss: 0.4459 - val_accuracy: 0.7154 - val_loss: 0.6137
 Epoch 56/100
 31/31 0s 11ms/step -
 accuracy: 0.7851 - loss: 0.4471 - val_accuracy: 0.7236 - val_loss: 0.6126
 Epoch 57/100
 31/31 0s 11ms/step -
 accuracy: 0.8389 - loss: 0.4017 - val_accuracy: 0.7154 - val_loss: 0.6160

Epoch 58/100
 31/31 0s 11ms/step -
 accuracy: 0.8213 - loss: 0.4298 - val_accuracy: 0.7073 - val_loss: 0.6192
 Epoch 59/100
 31/31 0s 12ms/step -
 accuracy: 0.7991 - loss: 0.4356 - val_accuracy: 0.7154 - val_loss: 0.6167
 Epoch 60/100
 31/31 0s 14ms/step -
 accuracy: 0.7764 - loss: 0.4654 - val_accuracy: 0.7236 - val_loss: 0.6188
 Epoch 61/100
 31/31 0s 11ms/step -
 accuracy: 0.8109 - loss: 0.4222 - val_accuracy: 0.7236 - val_loss: 0.6191
 Epoch 62/100
 31/31 0s 11ms/step -
 accuracy: 0.8115 - loss: 0.4322 - val_accuracy: 0.7236 - val_loss: 0.6234
 Epoch 63/100
 31/31 0s 11ms/step -
 accuracy: 0.8212 - loss: 0.4080 - val_accuracy: 0.7073 - val_loss: 0.6259
 Epoch 64/100
 31/31 0s 12ms/step -
 accuracy: 0.8328 - loss: 0.4136 - val_accuracy: 0.7236 - val_loss: 0.6225
 Epoch 65/100
 31/31 0s 11ms/step -
 accuracy: 0.8306 - loss: 0.4078 - val_accuracy: 0.6992 - val_loss: 0.6270
 Epoch 66/100
 31/31 0s 11ms/step -
 accuracy: 0.8211 - loss: 0.4183 - val_accuracy: 0.7236 - val_loss: 0.6230
 Epoch 67/100
 31/31 0s 11ms/step -
 accuracy: 0.8231 - loss: 0.4017 - val_accuracy: 0.7236 - val_loss: 0.6253
 Epoch 68/100
 31/31 0s 11ms/step -
 accuracy: 0.7960 - loss: 0.4539 - val_accuracy: 0.7236 - val_loss: 0.6241
 Epoch 69/100
 31/31 0s 11ms/step -
 accuracy: 0.8266 - loss: 0.4142 - val_accuracy: 0.6992 - val_loss: 0.6299
 Epoch 70/100
 31/31 0s 11ms/step -
 accuracy: 0.8028 - loss: 0.4208 - val_accuracy: 0.7236 - val_loss: 0.6302
 Epoch 71/100
 31/31 0s 11ms/step -
 accuracy: 0.7935 - loss: 0.4547 - val_accuracy: 0.7236 - val_loss: 0.6310
 Epoch 72/100
 31/31 0s 11ms/step -
 accuracy: 0.8172 - loss: 0.4122 - val_accuracy: 0.7154 - val_loss: 0.6309
 Epoch 73/100
 31/31 0s 11ms/step -
 accuracy: 0.8267 - loss: 0.3965 - val_accuracy: 0.7154 - val_loss: 0.6319

Epoch 74/100
 31/31 0s 11ms/step -
 accuracy: 0.8045 - loss: 0.4540 - val_accuracy: 0.7073 - val_loss: 0.6334
 Epoch 75/100
 31/31 0s 11ms/step -
 accuracy: 0.8018 - loss: 0.4327 - val_accuracy: 0.7154 - val_loss: 0.6319
 Epoch 76/100
 31/31 0s 12ms/step -
 accuracy: 0.8055 - loss: 0.4294 - val_accuracy: 0.7073 - val_loss: 0.6315
 Epoch 77/100
 31/31 0s 14ms/step -
 accuracy: 0.8105 - loss: 0.4329 - val_accuracy: 0.7073 - val_loss: 0.6358
 Epoch 78/100
 31/31 0s 13ms/step -
 accuracy: 0.7996 - loss: 0.4484 - val_accuracy: 0.7154 - val_loss: 0.6336
 Epoch 79/100
 31/31 0s 12ms/step -
 accuracy: 0.7919 - loss: 0.4350 - val_accuracy: 0.7154 - val_loss: 0.6392
 Epoch 80/100
 31/31 0s 11ms/step -
 accuracy: 0.8316 - loss: 0.4087 - val_accuracy: 0.7236 - val_loss: 0.6379
 Epoch 81/100
 31/31 0s 11ms/step -
 accuracy: 0.8195 - loss: 0.4149 - val_accuracy: 0.7073 - val_loss: 0.6406
 Epoch 82/100
 31/31 0s 11ms/step -
 accuracy: 0.8159 - loss: 0.4249 - val_accuracy: 0.7073 - val_loss: 0.6380
 Epoch 83/100
 31/31 0s 12ms/step -
 accuracy: 0.8179 - loss: 0.4314 - val_accuracy: 0.7154 - val_loss: 0.6381
 Epoch 84/100
 31/31 0s 11ms/step -
 accuracy: 0.7992 - loss: 0.4497 - val_accuracy: 0.7236 - val_loss: 0.6399
 Epoch 85/100
 31/31 0s 11ms/step -
 accuracy: 0.8168 - loss: 0.4189 - val_accuracy: 0.7073 - val_loss: 0.6409
 Epoch 86/100
 31/31 0s 12ms/step -
 accuracy: 0.8217 - loss: 0.4256 - val_accuracy: 0.7236 - val_loss: 0.6420
 Epoch 87/100
 31/31 0s 11ms/step -
 accuracy: 0.8184 - loss: 0.4322 - val_accuracy: 0.7154 - val_loss: 0.6443
 Epoch 88/100
 31/31 0s 12ms/step -
 accuracy: 0.8323 - loss: 0.4240 - val_accuracy: 0.7073 - val_loss: 0.6448
 Epoch 89/100
 31/31 0s 11ms/step -
 accuracy: 0.8012 - loss: 0.4487 - val_accuracy: 0.7236 - val_loss: 0.6424

```

Epoch 90/100
31/31          0s 11ms/step -
accuracy: 0.8314 - loss: 0.4042 - val_accuracy: 0.7154 - val_loss: 0.6420
Epoch 91/100
31/31          0s 11ms/step -
accuracy: 0.8482 - loss: 0.3900 - val_accuracy: 0.7073 - val_loss: 0.6437
Epoch 92/100
31/31          0s 11ms/step -
accuracy: 0.8512 - loss: 0.3928 - val_accuracy: 0.6992 - val_loss: 0.6441
Epoch 93/100
31/31          0s 11ms/step -
accuracy: 0.8342 - loss: 0.4150 - val_accuracy: 0.6992 - val_loss: 0.6478
Epoch 94/100
31/31          0s 12ms/step -
accuracy: 0.8241 - loss: 0.4259 - val_accuracy: 0.6911 - val_loss: 0.6473
Epoch 95/100
31/31          0s 11ms/step -
accuracy: 0.8175 - loss: 0.4272 - val_accuracy: 0.6992 - val_loss: 0.6486
Epoch 96/100
31/31          0s 13ms/step -
accuracy: 0.8077 - loss: 0.4512 - val_accuracy: 0.7073 - val_loss: 0.6455
Epoch 97/100
31/31          0s 11ms/step -
accuracy: 0.8320 - loss: 0.4116 - val_accuracy: 0.6992 - val_loss: 0.6483
Epoch 98/100
31/31          0s 11ms/step -
accuracy: 0.8142 - loss: 0.4541 - val_accuracy: 0.6992 - val_loss: 0.6526
Epoch 99/100
31/31          0s 12ms/step -
accuracy: 0.7873 - loss: 0.4529 - val_accuracy: 0.7154 - val_loss: 0.6491
Epoch 100/100
31/31          0s 11ms/step -
accuracy: 0.8501 - loss: 0.3974 - val_accuracy: 0.6911 - val_loss: 0.6543

```

```

[20]: # Evaluating the model
loss, accuracy = model.evaluate(X_test, y_test, verbose=0)

print(f"Test Loss: {loss:.4f}")
print(f"Model Accuracy: {accuracy * 100:.2f}%")

```

Test Loss: 0.7250

Model Accuracy: 73.38%

1.3.4 3. Provide a schematic representation of the neural network with its layers and activation functions. [1 Mark]

Below is the schematic representation of the network:

→ Input Layer (8 neurons)

Receives 8 input features.

-> Hidden Layer 1 (8 neurons, ReLU activation)

Applies the ReLU function to learn complex patterns.

-> Hidden Layer 2 (4 neurons, ReLU activation)

Further reduces dimensionality while preserving essential features.

-> Hidden Layer 3 (4 neurons, ReLU activation)

Extracts high-level patterns.

-> Output Layer (1 neuron, Sigmoid activation)

Outputs a probability value for binary classification.

```
[21]: # Importing required packages
import matplotlib.pyplot as plt
import matplotlib.patches as patches

def draw_neural_network(ax, left, right, bottom, top, layer_sizes, activations):
    # Calculate vertical and horizontal spacing
    v_spacing = (top - bottom) / float(max(layer_sizes))
    h_spacing = (right - left) / float(len(layer_sizes) - 1)

    # Draw nodes and activation functions
    for n, layer_size in enumerate(layer_sizes):
        layer_top = v_spacing * (layer_size - 1) / 2. + (top + bottom) / 2.

        for m in range(layer_size):
            # Create circles for nodes
            circle = plt.Circle((n * h_spacing + left, layer_top - m *
↪v_spacing), v_spacing / 4.,
                                color='w', ec='k', zorder=4)
            ax.add_artist(circle)

            # Annotate activation functions
            if n < len(layer_sizes) - 1 and m == layer_size - 1:
                ax.text(n * h_spacing + left, layer_top - m * v_spacing -
↪v_spacing / 2, activations[n], fontsize=12, ha='center')

    # Draw edges between layers
    for n, (layer_size_a, layer_size_b) in enumerate(zip(layer_sizes[:-1],
↪layer_sizes[1:])):
        layer_top_a = v_spacing * (layer_size_a - 1) / 2. + (top + bottom) / 2.
        layer_top_b = v_spacing * (layer_size_b - 1) / 2. + (top + bottom) / 2.

        for m in range(layer_size_a):
            for o in range(layer_size_b):
```

```

        # Draw lines for edges
        line = plt.Line2D([n * h_spacing + left, (n + 1) * h_spacing +
↪left],
                        [layer_top_a - m * v_spacing, layer_top_b - o
↪* v_spacing], c='k')
        ax.add_artist(line)

# Neural network architecture
layer_sizes = [8, 8, 4, 4, 1]
activations = ['ReLU', 'ReLU', 'ReLU', 'Sigmoid']

# Create figure and axis
fig, ax = plt.subplots(figsize=(12, 8))
ax.axis('on')

# Draw the neural network
draw_neural_network(ax, 0.1, 0.9, 0.1, 0.9, layer_sizes, activations)

# Annotate layer names
for i, layer_size in enumerate(layer_sizes):
    ax.text(i * (0.8 / (len(layer_sizes) - 1)) + 0.1, 0.95, f'Layer {i + 1}',
↪fontsize=16, ha='center')

plt.show()

```



